

Chat.AI — Enterprise Multi-Tenant AI SaaS Platform

Built on the GSD 14-Agent Autonomous Development Pipeline

Chat.AI is Technijian's upcoming multi-tenant AI SaaS product — and the reference implementation for every enterprise AI build Technijian delivers. Built end-to-end with GSD, the 14-agent autonomous pipeline that drives .NET + React + SQL Server from project description to alpha deployment.

CAPABILITY	STATUS	SCALE	PIPELINE
AI Dev · Full-Stack · Platform	Alpha live · GA in 4 weeks	~29,800 files, monorepo	14-agent GSD autonomous

WHAT TECHNIJIAN BUILT

A **production-grade multi-tenant AI SaaS platform** — not a prototype. Chat.AI is Technijian's own upcoming product launch, which doubles as the **reference implementation** for the custom AI work Technijian offers to enterprise clients. Same stack, same pipeline, same quality bar on your project.

- ✓ **Chat** — multi-LLM conversational UI with threaded history and per-thread model selection
- ✓ **Council** — same prompt across multiple LLMs concurrently, results compared, transcripts retained
- ✓ **Projects** — long-running workstreams with per-project memory, assistants, files, connectors
- ✓ **Assistants** — custom system-prompted agents, CRUD-managed, tenant-shareable
- ✓ **Connectors & Workflows** — outbound integrations (n8n, Microsoft Graph) and visual workflow builder with undo history

LLM BREADTH (NO LOCK-IN)

OpenAI

Anthropic

Google Vertex

Azure OpenAI

Cohere

Mistral

Groq

Z.ai

COMPLIANCE FROM DAY ONE

HIPAA audit trail

Tenant isolation (SESSION_CONTEXT)

GDPR data-subject rights

Rate limiting (user + tenant)

Redis security monitoring

Azure AD / MSAL + JWT

STACK

React 18 + TS + Vite

TailwindCSS + Zustand

ASP.NET Core 8

Dapper + SQL Server

SP-only data access

Figma design contract

THE 14-AGENT GSD PIPELINE

The interesting part isn't that Technijian built Chat.AI — it's **how fast, at what quality**. The same 14-agent autonomous pipeline is what Technijian runs on enterprise client projects.

PHASE A-E RequirementsAgent	PHASE A-E ArchitectureAgent	PHASE A-E FigmaIntegrationAgent	PHASE A-E PhaseReconcileAgent
PHASE A-E BlueprintFreezeAgent	PHASE A-E ContractFreezeAgent	PHASE F-G Orchestrator	PHASE F-G BlueprintAnalysisAgent
PHASE F-G CodeReviewAgent	PHASE F-G RemediationAgent	PHASE F-G QualityGateAgent	PHASE F-G E2EValidationAgent
PHASE F-G DeployAgent	PHASE F-G PostDeployValidationAgent		

BUILD OUTCOMES (CHATAI V8)

450 REQUIREMENTS REVIEWED	521 → 0 CRITICAL ISSUES	3x3 MODEL REVIEW ROUNDS	4wk LAUNCH TARGET
------------------------------	----------------------------	----------------------------	----------------------

Same pipeline we shipped Chat.AI with. Now running on your AI build.

949.379.8500

WHY THE PIPELINE WORKS IN PRODUCTION


Most autonomous dev systems don't survive first contact with a real codebase. GSD does, because four specific decisions are baked in from day one:

- ✓ **Vault-backed memory.** Every agent's system prompt, every architecture contract, every orchestrator decision lives in an Obsidian vault indexed by the agents themselves. Humans edit prompts in the vault, not in code — runtime picks up the change on next run.
- ✓ **Dual-auth routing for \$0 marginal cost.** Primary path uses CLI OAuth subscriptions (Claude Code, OpenAI Codex, Gemini). API keys are a backup path that only triggers on rate limit. Steady-state cost stays at the fixed subscription price.
- ✓ **Impact-analysis gates.** No agent renames a symbol without running a call-graph impact analysis first. HIGH / CRITICAL risk halts for human review. No blind refactors.
- ✓ **Three-model code review.** Every review cycle runs Claude, Codex, and Gemini in parallel against the same diff. Findings merged; RemediationAgent picks up the union set.

CODE-INTELLIGENCE AUGMENTATION STACK

- ✓ **Graphify** — knowledge graphs across the codebase
- ✓ **GitNexus** — impact analysis before every symbol edit
- ✓ **Context7** — live, up-to-date library documentation
- ✓ **Semgrep + OWASP** — security scanning
- ✓ **Playwright** — browser validation of the deployed app
- ✓ **GitHub MCP** — repo operations
- ✓ **Shannon** — code-review orchestration

WHAT THE ADMIN SURFACE INCLUDES



CAPABILITY	WHAT IT DELIVERS TO AN ENTERPRISE AI BUILD
Custom AI products	Not an AI skin on top of a SaaS tool — delivered with the same pipeline that shipped Chat.AI itself.
Compliance built in	HIPAA audit, tenant isolation, GDPR data-subject rights, and rate limiting are part of the generator, not a post-hoc patch.
Multi-LLM independence	Provider abstraction with 8 LLM backends (OpenAI, Anthropic, Google, Azure, Cohere, Mistral, Groq, Z.ai) — no single-vendor lock-in.
Design-first discipline	Figma spec is the contract, not a suggestion — FigmaIntegrationAgent validates the full 12-of-12 Figma Make deliverables.
Three-model review	Every code-review cycle runs Claude + Codex + Gemini in parallel against the same diff; findings merged for remediation.
Impact analysis	No agent renames a symbol without a call-graph impact analysis; HIGH / CRITICAL risk escalates to human review.
Evidence-backed	Chat.AI shipped with 450 requirements reviewed, 521 issues identified, 0 critical at verification — documented in the handoff, not marketing copy.

REPRESENTATIVE OUTCOME

“Technijian is launching its own multi-tenant AI SaaS — multi-LLM, HIPAA-compliant, tenant-isolated, audit-logged — built end-to-end with the 14-agent GSD pipeline. The same pipeline runs on your AI build: three-model code review, impact-analysis gates, Figma-as-contract, and deploy with automated rollback. Ask to see the alpha.”

— Flagship capability summary · based on Chat.AI v8 handoff evidence & GSD pipeline architecture