



ENTERPRISE AI PLATFORM

Technijian Chat.AI

Secure, multi-model AI with collective intelligence, workflow automation, and enterprise governance — purpose-built for organizations that demand more from AI.

4+

AI PROVIDERS

32+

ENDPOINTS

>60%

COUNCIL WIN RATE

99.9%

UPTIME SLA

The Challenge

Enterprises adopting AI face critical obstacles that generic chatbot subscriptions cannot solve.

Vendor Lock-In

Tied to a single AI provider with no failover, no model comparison, and no negotiating leverage. When that provider has an outage or raises prices, your business stalls.



Chat.AI connects 4+ providers with automatic failover

Runaway AI Costs

No visibility into AI spend. Bills arrive as surprises. You can't attribute costs to the projects or departments actually driving the work, and no platform stops the bleed before it happens.



Credit metering with hard spend caps & real-time alerts

Security & Compliance Gaps

Consumer AI tools lack audit trails, role-based access, and tenant isolation. Legal and compliance teams cannot approve them for regulated data or sensitive workflows.



Row-level isolation, full audit trail, HIPAA/SOC 2 ready

Single-Model Hallucinations

One model means no cross-validation. Wrong answers go unchecked and unverified, putting your team at risk when AI output drives real business decisions.



AI Council cross-validates with 3+ models per query

4+

AI PROVIDERS UNIFIED

>60%

AI COUNCIL WIN RATE

>30%

HALLUCINATION
REDUCTION

100%

AUDIT TRAIL COVERAGE

Core Capabilities

Six integrated modules that transform AI from a siloed chatbot into an enterprise intelligence platform.



Multi-Model AI Engine

Access GPT-5, Claude Sonnet 4.6, Gemini 2.5 Pro, Nova, Mistral — all in one platform.

- ✓ Switch models mid-conversation
- ✓ Smart routing across Economy / Standard / Frontier
- ✓ Streaming with file attachments
- ✓ Thread sharing & full history



AI Council (Patent Pending)

Multi-model collective intelligence that outperforms any single model.

- ✓ 3-stage: Diverge → Converge → Synthesize
- ✓ Anonymized peer review across models
- ✓ >60% win rate vs single models
- ✓ >30% hallucination reduction



Knowledge & RAG

Enterprise knowledge base with vector search across your documents.

- ✓ Qdrant, Weaviate, Milvus, pgvector, OpenSearch
- ✓ File vectorization & semantic search
- ✓ SharePoint & Google Drive sync
- ✓ Auto-upgrade keeps KB working at scale



Custom Assistants

Build purpose-specific AI agents with custom system prompts and connectors.

- ✓ System prompts & personas
- ✓ REST, GraphQL, DB connectors
- ✓ Project workspaces with context
- ✓ Template marketplace



Workflow Automation

Trigger complex workflows from chat context, powered by N8N integration.

- ✓ N8N workflow orchestration
- ✓ Connect to any API endpoint
- ✓ Scheduled & event-driven flows
- ✓ Chat-triggered automations



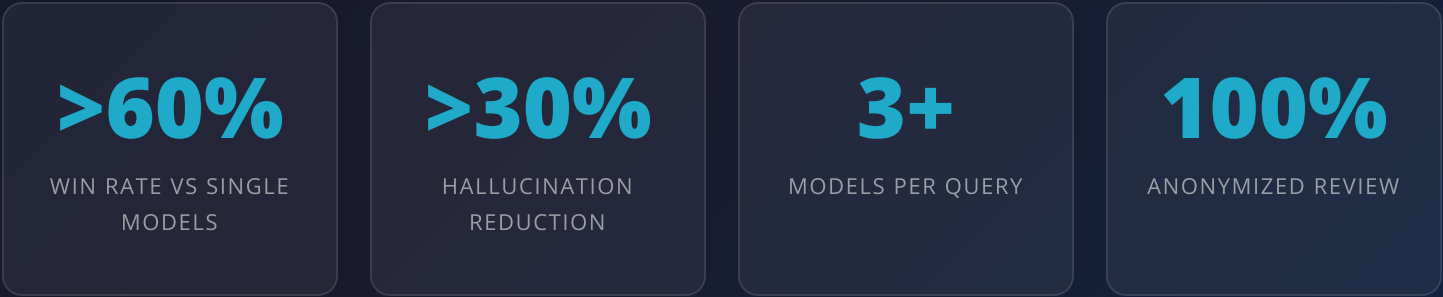
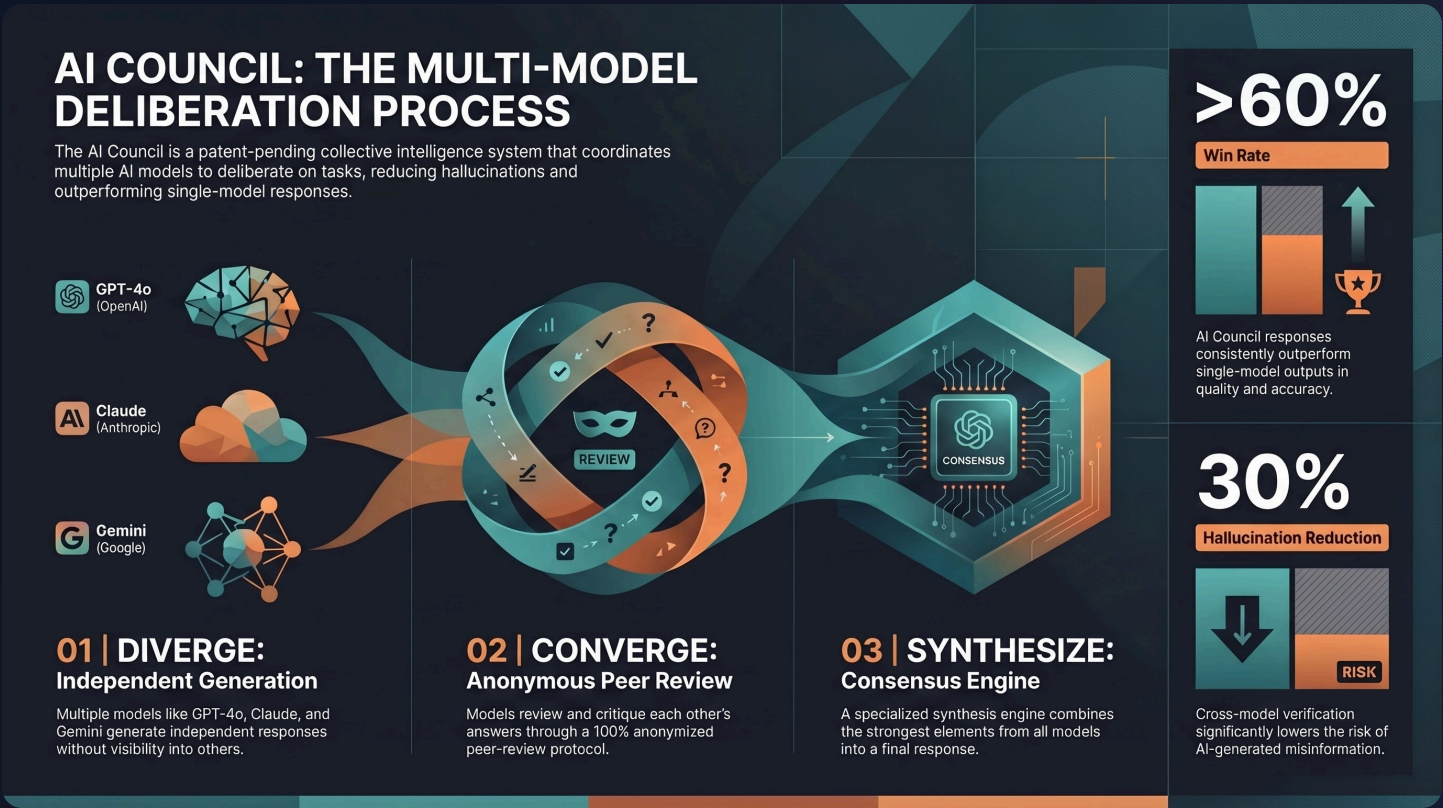
Enterprise Governance

Complete visibility over AI usage, costs, and compliance.

- ✓ Credit-level cost attribution
- ✓ Hard spend caps per tenant
- ✓ Real-time usage dashboard
- ✓ 80/90/95/100% alerts & warn-upgrade

The AI Council — Patent Pending

Our multi-model deliberation system synthesizes collective intelligence, achieving consensus that outperforms any single model.



Real-World Use Cases

How organizations across industries leverage Chat.AI to transform their AI operations.

Financial Services

Compliance-First AI Deployment

A 200-person financial firm needed AI for research but couldn't risk data leaks or hallucinated financial advice.

Solution: On-premise deployment with Azure AD SSO, full audit trails, and AI Council cross-validation to eliminate hallucinations on financial questions.

- ✓ 50 users across research, compliance, risk
- ✓ FINRA compliance review passed
- ✓ AI Council flagged 3 hallucinations/week
- ✓ 40% faster research turnaround

\$24,000

ANNUAL SAVINGS (50 USERS)

Healthcare Organization

HIPAA-Compliant AI Knowledge Base

Clinical staff needed AI-assisted search across protocols with full HIPAA compliance and audit trails.

Solution: RAG knowledge base with 500+ clinical protocols, row-level tenant isolation, and full audit logging of every query.

- ✓ 120 clinical users onboarded
- ✓ Zero PHI data leakage incidents
- ✓ HIPAA Security Rule compliant
- ✓ 3-min average query-to-answer

100%

AUDIT TRAIL COVERAGE

Technology Company

Multi-Model Developer Platform

100+ developers needed multiple AI models for code review and documentation without separate subscriptions.

Solution: Unified access to GPT-5, Claude Sonnet, Gemini 2.5 Pro, and Nova through a single platform with credit spend caps per workspace.

- ✓ 250 developers, all models on tap
- ✓ Consolidated 3 AI subscriptions
- ✓ Smart routing — Economy for code, Frontier for review
- ✓ 60% faster code review cycles

\$168K

ANNUAL SAVINGS (250 USERS)

Editions & Pricing

Pay for usage, not seats. Three hosted editions plus on-premise. Every edition includes every public AI model. Annual prepay saves 17%.

Team

\$750/mo

\$7,500/yr · save 17%

- ✓ Unlimited users
- ✓ 30M credits/month included
- ✓ 1 TB storage
- ✓ 3 MCP servers · 10 Assistants
- ✓ 1 shared knowledge base
- ✓ All public AI providers
- ✓ Email 48hr · 99.5% SLA
- ✓ 90-day audit retention

MOST POPULAR

Business

\$3,500/mo

\$35,000/yr · save 17%

- ✓ Unlimited users
- ✓ 300M credits/month included
- ✓ 5 TB storage
- ✓ 10 MCP servers · unlimited Assistants
- ✓ Multi-tenant admin · white-label
- ✓ SSO/SAML + SCIM + audit logs
- ✓ Priority 4hr · 99.9% SLA
- ✓ 1-year audit retention

Enterprise / Compliant

Talk to us

Typically \$9,000+/mo

- ✓ Unlimited users
- ✓ 1B+ credits/month or custom
- ✓ Custom storage · dedicated KB
- ✓ Unlimited MCP servers
- ✓ BAA · HIPAA · SOC 2 Type II
- ✓ Always Encrypted · custom DPA
- ✓ Named CSM 1hr · 99.95% SLA
- ✓ 7-year audit retention

On-Premise

~\$60K/yr+

Team-eq. · Business
~\$120K–\$190K

- ✓ Your cloud / datacenter / air-gap
- ✓ BYO LLM keys — \$2/1M orch fee
- ✓ BYO vector DB · BYO storage
- ✓ BYO MCP integrations
- ✓ BAA for Technijian components
- ✓ Quarterly release cadence
- ✓ Tiered support included
- ✓ Enterprise tier ~\$250K+/yr

No Seat Taxes

Pricing is by edition and usage — never per user. Roll out organization-wide without a seat-multiplier surprise.

Founding Customer Program

Pre-launch cohort gets rates locked for 24–36 months in exchange for being a named reference. Multi-year adds another 5–10%.

Never Broken

Reads always work. If you outgrow a limit, the platform warns early and (with your permission) auto-upgrades. No silent cut-offs.

Credit Meter & Resources

AI usage is metered in credits — 1 credit equals 1 token on a standard-class model. Commit and save; bring your own and save more.

Commit & Save

Monthly Commitment	\$ / 1M Credits
Pay-as-you-go	\$18
\$2,500 / mo	\$15
\$10,000 / mo	\$13
\$40,000 / mo	\$11
\$100,000+ / mo	\$9

Frontier-class consumption floors at \$30 per 1M tokens effective. Hard spend caps on by default — no surprise invoices.

Model Classes & Modifiers

Class	Examples	/ token
Economy	Haiku, GPT-5 mini, Nova Lite	0.5×
Standard	Sonnet, GPT-5, Gemini 2.5 Pro	1.0×
Frontier	Opus, GPT-5.4, o3, Gemini 3 Pro	2.5×
Cached input	Repeated context	0.25×
Batch / async	Bulk & overnight jobs	0.6×

Route simple queries to Economy, reasoning to Frontier — mix & match per assistant.

Hosted Resources (Beyond Included Allowance)

Dedicated Knowledge Base

Isolated vector DB — Qdrant, Weaviate, Milvus, pgvector, OpenSearch.

Small · ~10M vectors · \$500/mo
Medium · ~50M vectors · \$1,200/mo
Large · ~200M vectors · \$2,500/mo
X-Large · 200M+ · \$5,000/mo

Auto-upgrade keeps your KB working.

Storage

Documents, embeddings, history, audit logs, exports.

\$50/TB/mo

Billed in 1 TB blocks. Team includes 1 TB, Business 5 TB, Enterprise custom.

BYO storage on On-Premise = free.

MCP Servers

Tool integrations assistants & workflows call.

Public catalog · \$50–\$75/mo
Technician-built · \$150–\$250/mo
Bespoke custom · \$8K–\$25K + \$150–\$500/mo
Bring your own · Free on hosted

Build once, sell many — design-partner credit.

Competitive Comparison

Chat.AI delivers enterprise capabilities that consumer AI tools cannot match.

CAPABILITY	CHATGPT TEAM	MICROSOFT COPILOT	GOOGLE GEMINI	TECHNIJIAN CHAT.AI
Multi-Model Access	—	—	—	✓ 4+ Providers
AI Council Consensus	—	—	—	✓ Patent Pending
Credit-Level Cost Control	—	—	—	✓ Spend Caps
On-Premise Deployment	—	—	—	✓ Full Stack
Complete Audit Trail	—	Limited	Limited	✓ Full Trail
Workflow Automation	—	—	—	✓ N8N
RAG Knowledge Base	—	Basic	—	✓ Vector Search
Custom Assistants	GPTs (Limited)	—	Gems (Limited)	✓ Full Platform
Tenant Isolation	—	Partial	—	✓ Row-Level
Starting price (50 users)	\$1,250	\$1,500	\$1,000	\$750 flat

Why Chat.AI Outperforms

Consumer AI tools like ChatGPT Team, Copilot, and Gemini are optimized for individual productivity — not enterprise operations. They lack the governance, observability, and architectural flexibility that regulated industries and large organizations require.

Chat.AI is purpose-built for enterprise AI deployments: multi-provider access, patent-pending AI Council validation, full audit trails, and deep integrations with your existing tech stack.

\$3,500

Monthly savings vs
3 fragmented AI tools

100-USER TEAM · \$42K/YR

Return on Investment

Chat.AI delivers measurable savings from day one by consolidating multiple AI subscriptions.

SMALL TEAM • 25 USERS

\$1,375/mo (ChatGPT+
Copilot)

\$750/mo — Chat.AI Team

\$7,500

ANNUAL SAVINGS

MID-MARKET • 100 USERS

\$7,000/mo (3 fragmented
tools)

\$3,500/mo — Chat.AI
Business

\$42,000

ANNUAL SAVINGS

ENTERPRISE • 500 USERS

\$35,000/mo (3 fragmented
tools)

\$9,000/mo — Chat.AI
Enterprise

\$312,000

ANNUAL SAVINGS

Beyond Cost Savings

- ✓ **Productivity gains:** 40% faster research & writing
- ✓ **Quality gains:** >30% fewer hallucinations via AI Council
- ✓ **Risk reduction:** audit-ready for HIPAA, SOC 2, PCI
- ✓ **Vendor flexibility:** no lock-in, automatic failover

Payback Period

< 90

DAYS FOR 50+ USER ORGS

Most enterprise customers recover their Chat.AI investment in under 90 days through AI subscription consolidation alone — then compound savings across year 1 and beyond.

Security & Compliance

Enterprise-grade security meeting the most stringent compliance requirements.

**Azure AD SSO**
Single sign-on with Microsoft identity. Zero friction.

**Row-Level Security**
Complete tenant isolation, no data leakage.

**Full Audit Trail**
Every interaction logged with timestamps.

**Azure Key Vault**
Secrets managed with automatic rotation.

**Prompt Injection Defense**
Multi-layer input validation and sanitization.

**HIPAA / SOC 2 Ready**
Healthcare and financial compliance.

**RBAC Roles**
Granular role-based access control.

**Rate Limiting**
Per-user, per-model endpoint controls.

Compliance Frameworks Supported


HIPAA


SOC 2


PCI-DSS


GDPR


NIST CSF


ISO 27001

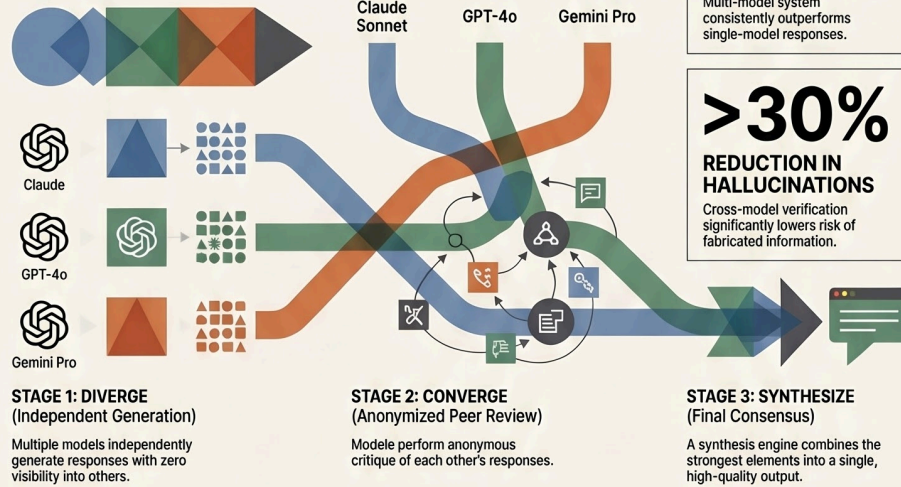
Chat.AI’s architecture is built from the ground up with compliance in mind. Enterprise customers receive full audit support and can request SOC 2 Type II reports.

Chat.AI Intelligence Layer


Enterprise multi-model AI architecture — how the AI Council delivers superior results

Technijian Chat.AI: The Enterprise Multi-Model Intelligence Layer


THE AI COUNCIL DELIBERATION ENGINE



ADVANCED ARCHITECTURE & KNOWLEDGE



ENTERPRISE FOUNDATION & METRICS





Enterprise AI, Reimagined.

technijian.com

Schedule a demo: RJain@technijian.com • 949.379.8500

