

Secure, multi-model AI with collective intelligence, workflow automation, and enterprise governance — all in one platform.

KEY CAPABILITIES

Multi-Model AI

- ✓ GPT-5, Claude Sonnet 4.6, Gemini 2.5 Pro, Nova, Mistral & more
- ✓ Switch models mid-conversation or auto-route by task
- ✓ Streaming responses with file attachments & code analysis
- ✓ Thread sharing with permissions, export & full history

AI Council (Patent Pending)

- ✓ 3-stage deliberation: Diverge → Converge → Synthesize
- ✓ 3+ models per query with anonymized peer review
- ✓ >60% win rate vs single models across benchmarks
- ✓ >30% hallucination reduction through cross-validation

Custom Assistants & Projects

- ✓ Purpose-built AI assistants with custom system prompts
- ✓ REST, GraphQL, and direct database connector integrations
- ✓ Project workspaces with inherited context & team sharing
- ✓ Template marketplace for rapid assistant deployment

Workflow & Knowledge

- ✓ N8N workflow automation triggered from chat context
- ✓ RAG with Weaviate, Pinecone, or Chroma vector stores
- ✓ File vectorization, semantic search & persistent org memory
- ✓ Knowledge base auto-sync with SharePoint & Google Drive

HOW IT WORKS

- 1 Deploy** — Provision your tenant with Azure AD SSO in under 15 minutes
- 2 Configure** — Pick models, set credit spend caps, build assistants & knowledge bases
- 3 Scale** — Add users, workflows, knowledge bases & integrations as you grow

VS. COMPETITORS

Capability	ChatGPT	Copilot	Gemini	Chat.AI
Multi-Model Access	—	—	—	✓ 4+ Providers
AI Council Consensus	—	—	—	✓ Patent Pending
Credit-Level Cost Control	—	—	—	✓ Spend Caps
On-Premise Deployment	—	—	—	✓ Full Stack
Complete Audit Trail	—	Limited	Limited	✓ Full
Workflow Automation	—	—	—	✓ N8N
Starting price (50 users)	\$1,250	\$1,500	\$1,000	\$750 flat

WHY CHAT.AI

- **Multi-Provider Freedom:** No vendor lock-in — OpenAI, Anthropic, Google, Azure, and open-source with automatic failover
- **No Surprise Invoices:** Hard spend caps on by default, real-time credit alerts, no per-seat taxes — pay only for what you use
- **BAA-Backed Compliance:** Routing locked to Amazon Bedrock + Azure OpenAI for Compliant edition; HIPAA, SOC 2 Type II, GDPR
- **On-Premise & BYO:** Run in your datacenter, bring your own LLM keys (\$2/1M orch fee), vector DB, storage — no lock-in
- **Council Consensus:** Patent-pending multi-agent deliberation achieves >60% win rate & >30% fewer hallucinations
- **Workflow Automation:** Trigger N8N workflows from chat, connect to any API, and build custom AI-powered pipelines

EDITIONS & PRICING

Edition	From	Credits/mo	Storage	SLA
Team	\$750/mo	30M	1 TB	99.5%
Business	\$3,500/mo	300M	5 TB	99.9%
Enterprise / Compliant	Talk to us	1B+	Custom	99.95%
On-Premise	~\$60K/yr	BYO keys	BYO	BAA

Annual saves 17% • Unlimited users on every edition • Credits: \$18/1M PAYG → \$9/1M at commit

SECURITY & COMPLIANCE

- ✓ Azure AD SSO
- ✓ Row-Level Security
- ✓ Full Audit Trail
- ✓ Azure Key Vault
- ✓ Prompt Injection Defense
- ✓ HIPAA / SOC 2 Ready
- ✓ RBAC Role Management
- ✓ Rate Limiting & Throttling

WAYS TO SAVE ON CREDITS

COMMIT VOLUME
\$18 → \$9 per 1M

50% off

SMART ROUTING
Standard → Economy

50% off

CACHED INPUT
Repeated context

75% off