TECHNIJIAN
technology as a solution

ENTERPRISE AI PLATFORM

Technijian Chat.AI

Secure, multi-model AI with collective intelligence, workflow automation, and enterprise governance — purpose-built for organizations that demand more from AI.

4+	32+	>60%	99.9%
AI PROVIDERS	ENDPOINTS	COUNCIL WIN RATE	UPTIME SLA

The Challenge

Enterprises adopting AI face critical obstacles that generic chatbot subscriptions cannot solve.

Vendor Lock-In

Tied to a single AI provider with no failover, no model comparison, and no negotiating leverage. When that provider has an outage or raises prices, your business stalls.

→ **Chat.AI connects 4+ providers with automatic failover**

Runaway AI Costs

No visibility into which teams or users consume the most tokens. Bills arrive as surprises, and you can't attribute AI costs to the projects or departments actually driving the spend.

→ **Token-level metering with per-user budgets & alerts**

Security & Compliance Gaps

Consumer AI tools lack audit trails, role-based access, and tenant isolation. Legal and compliance teams cannot approve them for regulated data or sensitive workflows.

→ **Row-level isolation, full audit trail, HIPAA/SOC 2 ready**

Single-Model Hallucinations

One model means no cross-validation. Wrong answers go unchecked and unverified, putting your team at risk when AI output drives real business decisions.

→ **AI Council cross-validates with 3+ models per query**

4+	>60%	>30%	100%
AI PROVIDERS UNIFIED	AI COUNCIL WIN RATE	HALLUCINATION REDUCTION	AUDIT TRAIL COVERAGE

Core Capabilities

Six integrated modules that transform AI from a siloed chatbot into an enterprise intelligence platform.

Multi-Model AI Engine

Access GPT-4o, Claude Sonnet, Gemini Pro, Llama through a unified interface.

- ✓ Switch models mid-conversation
- ✓ Auto-route queries by task type
- ✓ Streaming with file attachments
- ✓ Thread sharing & full history

AI Council (Patent Pending)

Multi-model collective intelligence that outperforms any single model.

- ✓ 3-stage: Diverge → Converge → Synthesize
- ✓ Anonymized peer review across models
- ✓ >60% win rate vs single models
- ✓ >30% hallucination reduction

Knowledge & RAG

Enterprise knowledge base with vector search across your documents.

- ✓ Weaviate, Pinecone, Chroma support
- ✓ File vectorization & semantic search
- ✓ SharePoint & Google Drive sync
- ✓ Persistent organizational memory

Custom Assistants

Build purpose-specific AI agents with custom system prompts and connectors.

- ✓ System prompts & personas
- ✓ REST, GraphQL, DB connectors
- ✓ Project workspaces with context
- ✓ Template marketplace

Workflow Automation

Trigger complex workflows from chat context, powered by N8N integration.

- ✓ N8N workflow orchestration
- ✓ Connect to any API endpoint
- ✓ Scheduled & event-driven flows
- ✓ Chat-triggered automations

Enterprise Governance

Complete visibility over AI usage, costs, and compliance.

- ✓ Token-level cost attribution
- ✓ Per-user & per-model budgets
- ✓ Usage analytics dashboard
- ✓ 75%/100% budget alerts

The AI Council — Patent Pending

Our multi-model deliberation system synthesizes collective intelligence, achieving consensus that outperforms any single model.

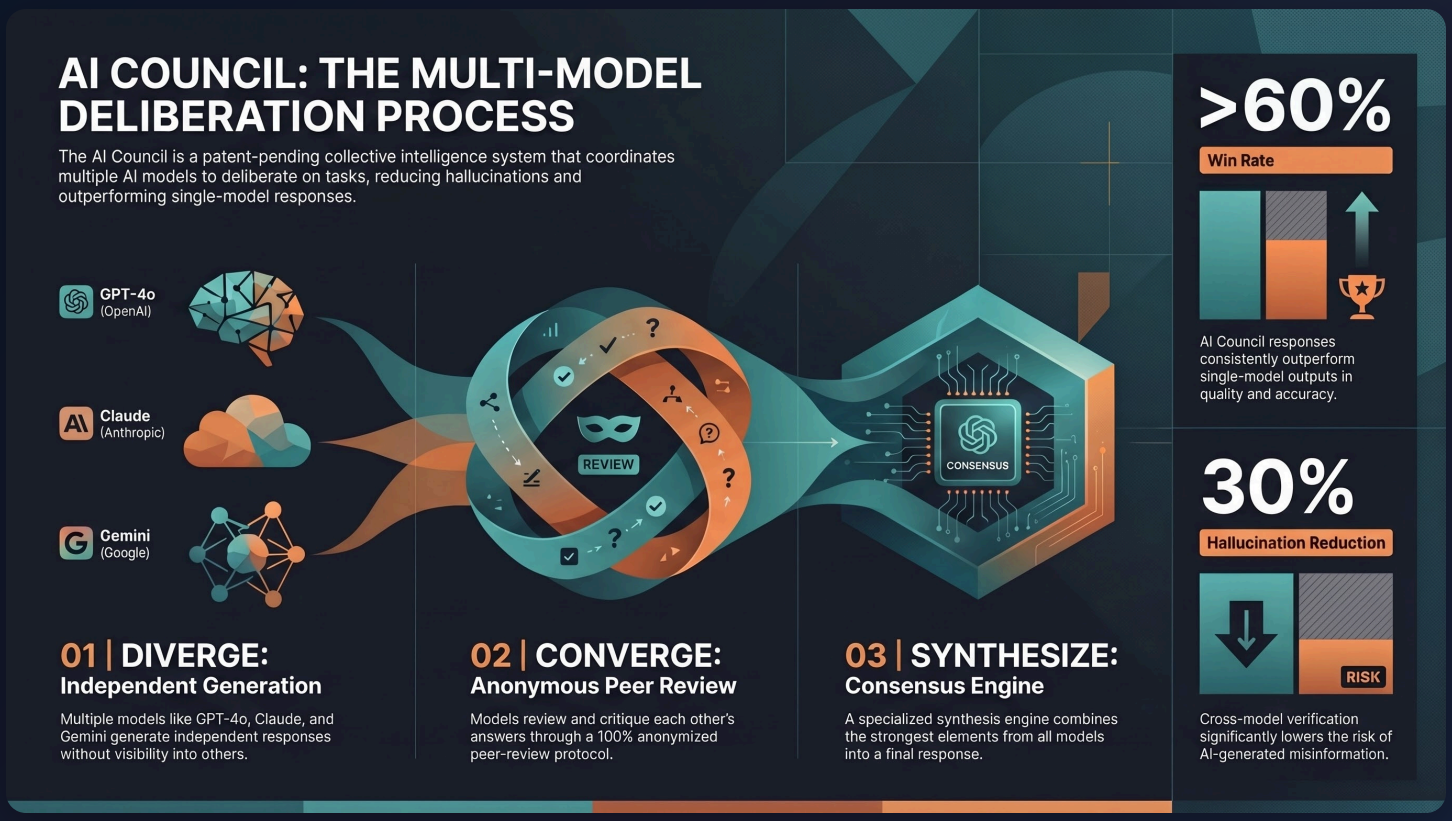

>60%	>30%	3+	100%
WIN RATE VS SINGLE MODELS	HALLUCINATION REDUCTION	MODELS PER QUERY	ANONYMIZED REVIEW

Real-World Use Cases

How organizations across industries leverage Chat.AI to transform their AI operations.

💼 Financial Services

Compliance-First AI Deployment

A 200-person financial firm needed AI for research but couldn't risk data leaks or hallucinated financial advice.

Solution: On-premise deployment with Azure AD SSO, full audit trails, and AI Council cross-validation to eliminate hallucinations on financial questions.

✓ 50 users across research, compliance, risk
✓ FINRA compliance review passed
✓ AI Council flagged 3 hallucinations/week
✓ 40% faster research turnaround

$19,212
ANNUAL SAVINGS (50 USERS)

🏥 Healthcare Organization

HIPAA-Compliant AI Knowledge Base

Clinical staff needed AI-assisted search across protocols with full HIPAA compliance and audit trails.

Solution: RAG knowledge base with 500+ clinical protocols, row-level tenant isolation, and full audit logging of every query.

✓ 120 clinical users onboarded
✓ Zero PHI data leakage incidents
✓ HIPAA Security Rule compliant
✓ 3-min average query-to-answer

100%
AUDIT TRAIL COVERAGE

🛠 Technology Company

Multi-Model Developer Platform

100+ developers needed multiple AI models for code review and documentation without separate subscriptions.

Solution: Unified access to GPT-4o, Claude Sonnet, Gemini Pro, and Llama through a single platform with per-user token budgets.

✓ 250 developers, 4+ AI models each
✓ Consolidated 5 AI subscriptions
✓ Per-developer token budgets
✓ 60% faster code review cycles

$174K
ANNUAL SAVINGS (250 USERS)

Plans & Pricing

Simple, transparent pricing. No per-seat charges. Annual plans save 17%.

Starter

$299

$2,990/yr

- ✓ Up to 10 users
- ✓ 10M tokens/month
- ✓ Multi-model access
- ✓ Azure AD SSO
- ✓ Basic audit trail
- ✓ Email support

Growth

$899

$8,990/yr

- ✓ Up to 50 users
- ✓ 35M tokens/month
- ✓ Custom assistants
- ✓ RAG knowledge base
- ✓ Full audit trail
- ✓ Priority support

MOST POPULAR

Professional

$1,699

$16,990/yr

- ✓ Up to 100 users
- ✓ 75M tokens/month
- ✓ AI Council included
- ✓ N8N automation
- ✓ Advanced analytics
- ✓ Dedicated CSM

Enterprise

Custom

Contact for pricing

- ✓ 250+ users
- ✓ 200M+ tokens/month
- ✓ On-premise ($45K/yr)
- ✓ Custom integrations
- ✓ SLA guarantees
- ✓ 24/7 phone support

No Per-Seat Charges

Pricing is based on usage (tokens), not seats. Add users freely without cost surprises.

Annual Plans Save 17%

Commit annually for a 17% discount. Most teams see ROI in under 90 days from subscription consolidation.

30-Day Pilot Program

Try Chat.AI risk-free with a 30-day pilot for qualifying organizations. Full features, full support.

Competitive Comparison

Chat.AI delivers enterprise capabilities that consumer AI tools cannot match.

CAPABILITY	CHATGPT TEAM	MICROSOFT COPILOT	GOOGLE GEMINI	TECHNIJIAN CHAT.AI
Multi-Model Access	—	—	—	✓ 4+ Providers
AI Council Consensus	—	—	—	✓ Patent Pending
Per-User Cost Tracking	—	—	—	✓ Token-Level
On-Premise Deployment	—	—	—	✓ Full Stack
Complete Audit Trail	—	Limited	Limited	✓ Full Trail
Workflow Automation	—	—	—	✓ N8N
RAG Knowledge Base	—	Basic	—	✓ Vector Search
Custom Assistants	GPTs (Limited)	—	Gems (Limited)	✓ Full Platform
Tenant Isolation	—	Partial	—	✓ Row-Level
Price per 50 Users/mo	$1,250	$1,500	$1,000	**$899**

Why Chat.AI Outperforms

Consumer AI tools like ChatGPT Team, Copilot, and Gemini are optimized for individual productivity — not enterprise operations. They lack the governance, observability, and architectural flexibility that regulated industries and large organizations require.

Chat.AI is purpose-built for enterprise AI deployments: multi-provider access, patent-pending AI Council validation, full audit trails, and deep integrations with your existing tech stack.

$351

Average monthly savings vs nearest competitor

BASED ON 50-USER PRICING

Return on Investment

Chat.AI delivers measurable savings from day one by consolidating multiple AI subscriptions.

SMALL TEAM • 10 USERS

~~$389/mo individual subscriptions~~

$299/mo with Chat.AI

$1,080

ANNUAL SAVINGS

MID-MARKET • 50 USERS

~~$2,500/mo individual subscriptions~~

$899/mo with Chat.AI

$19,212

ANNUAL SAVINGS

ENTERPRISE • 250 USERS

~~$18,500/mo individual subscriptions~~

$3,999/mo with Chat.AI

$174,012

ANNUAL SAVINGS

Beyond Cost Savings

- ✓ **Productivity gains:** 40% faster research & writing
- ✓ **Quality gains:** >30% fewer hallucinations via AI Council
- ✓ **Risk reduction:** audit-ready for HIPAA, SOC 2, PCI
- ✓ **Vendor flexibility:** no lock-in, automatic failover

Payback Period

< 90

DAYS FOR 50+ USER ORGS

Most enterprise customers recover their Chat.AI investment in under 90 days through AI subscription consolidation alone — then compound savings across year 1 and beyond.

Security & Compliance

Enterprise-grade security meeting the most stringent compliance requirements.

Azure AD SSO
Single sign-on with Microsoft identity. Zero friction.

Row-Level Security
Complete tenant isolation, no data leakage.

Full Audit Trail
Every interaction logged with timestamps.

Azure Key Vault
Secrets managed with automatic rotation.

Prompt Injection Defense
Multi-layer input validation and sanitization.

HIPAA / SOC 2 Ready
Healthcare and financial compliance.

RBAC Roles
Granular role-based access control.

Rate Limiting
Per-user, per-model endpoint controls.

Compliance Frameworks Supported

- ✓ HIPAA
- ✓ SOC 2
- ✓ PCI-DSS
- ✓ GDPR
- ✓ NIST CSF
- ✓ ISO 27001

Chat.AI's architecture is built from the ground up with compliance in mind. Enterprise customers receive full audit support and can request SOC 2 Type II reports.

Chat.AI Intelligence Layer

Enterprise multi-model AI architecture — how the AI Council delivers superior results

Technijian Chat.AI: The Enterprise Multi-Model Intelligence Layer

THE AI COUNCIL DELIBERATION ENGINE

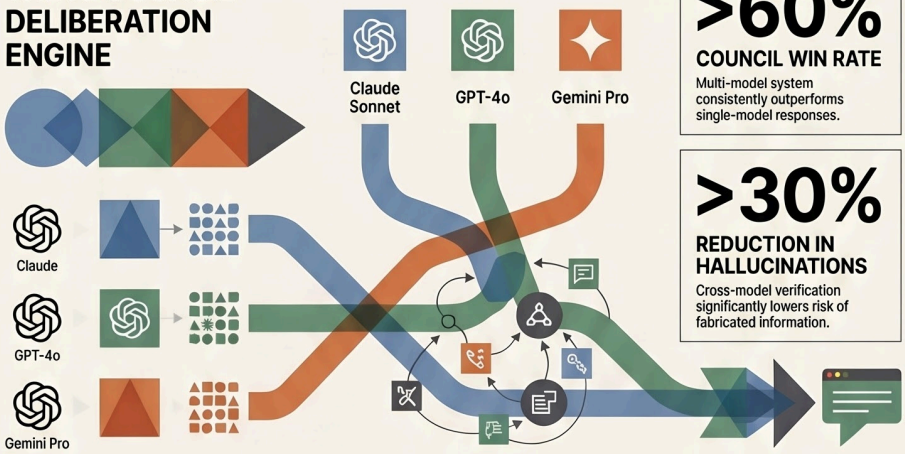

>60%
COUNCIL WIN RATE
Multi-model system consistently outperforms single-model responses.

>30%
REDUCTION IN HALLUCINATIONS
Cross-model verification significantly lowers risk of fabricated information.

STAGE 1: DIVERGE
(Independent Generation)
Multiple models independently generate responses with zero visibility into others.

STAGE 2: CONVERGE
(Anonymized Peer Review)
Modele perform anonymous critique of each other's responses.

STAGE 3: SYNTHESIZE
(Final Consensus)
A synthesis engine combines the strongest elements into a single, high-quality output.

ADVANCED ARCHITECTURE & KNOWLEDGE

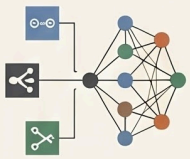

32+
API ENDPOINTS ACROSS 4+ PROVIDERS
Unified access to OpenAI, Anthropic, Google, and Meta models with mid-conversation switching.

MULTI-AGENT ORCHESTRATION

Advanced workflow automation from chat context, utilicing NBN for complex task sequencing.

RAG KNOWLEDGE SYSTEM
Persistent organizational memory using file vectorization and semantic search via Weaviate, Pinecone, or Chroma.

AUTO-SYNC ECOSYSTEM
Seamless knowledge base synchronization with major platforms.

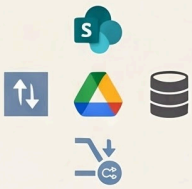

ENTERPRISE FOUNDATION & METRICS

99.9%
UPTIME SLA

Mission-critical reliability for high-availability environments.

4 FLEXIBLE DEPLOYMENT OPTIONS

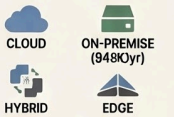

Meet specific data residency needs.

SECURITY & GOVERNANCE

Features Asure AC SSO, RBAC, fall audit trails, and Azure Key Vault encryption.

TOKEN-LEVEL COST ATTRIBUTION

Granaler bacaing per user and model, with automated 75% and 100% alerts.

COST EFFICIENCY COMPARISON (50 USERS)

PLATFORM	MONTHLY COST (50 USERS)	ENTERPRISE FEATURES
Technijian Chat.AI	$899	Ai Council, On-Prem, NBN, RAG
Google Gemini	$1,000	Standard AI
ChatGPT Team	$1,250	Standard AI
Microsoft Copilot	$1,500	Standard AI

TECHNIJIAN

technology as a solution

Enterprise AI, Reimagined.

technijian.com

Schedule a demo: RJain@technijian.com • 949.379.8500

© 2026 Technijian · 18 Technology Dr. Ste 141, Irvine CA 92618 · technijian.com